

Natural Attention: Theoretical Foundations

Jeronimo Aranda

December 17, 2024

Abstract

We present rigorous proofs for fundamental properties of natural attention mechanisms, focusing on energy conservation, scale invariance, and information processing. We clearly distinguish between proven results and conjectures requiring further development.

1 Preliminaries

Definition 1 (Natural Attention). For input sequence $X \in \mathbb{R}^{L \times d}$, the natural attention mechanism is defined as:

$$A(X) = \text{softmax} \left(\frac{XW_Q W_K^T X^T}{\sqrt{d}} \right) XW_V$$

where $W_Q, W_K, W_V \in \mathbb{R}^{d \times d}$ are learnable parameters.

Lemma 2 (Attention Energy Form). *The attention energy matrix E for a given head can be expressed as:*

$$E = \frac{QK^T}{\sqrt{d}} = \frac{XW_Q W_K^T X^T}{\sqrt{d}}$$

where $Q = XW_Q$ and $K = XW_K$ are query and key projections.

2 Energy Conservation

Theorem 3 (Energy Conservation). *For an attention mechanism with H heads and head dimension d , assuming normalized weight matrices, the total attention energy is bounded:*

$$\sum_{h=1}^H \text{tr}(E_h E_h^T) \leq d \cdot H \cdot \|X\|_F^4$$

where E_h is the energy matrix for head h .

Proof. First, consider a single head h :

$$E_h = \frac{XW_Q^{(h)}W_K^{(h)T}X^T}{\sqrt{d}}$$

By submultiplicativity of Frobenius norm:

$$\begin{aligned}\|E_h\|_F^2 &= \text{tr}(E_h E_h^T) \\ &= \frac{1}{d} \text{tr}(XW_Q^{(h)}W_K^{(h)T}X^T XW_K^{(h)}W_Q^{(h)T}X^T) \\ &\leq \frac{1}{d} \|X\|_F^2 \|W_Q^{(h)}\|_F^2 \|W_K^{(h)}\|_F^2 \|X\|_F^2\end{aligned}$$

For normalized weight matrices (standard initialization):

$$\|W_Q^{(h)}\|_F^2 = \|W_K^{(h)}\|_F^2 = d$$

Therefore:

$$\text{tr}(E_h E_h^T) \leq d \|X\|_F^4$$

Summing over all heads completes the proof. \square

3 Scale Invariance

Theorem 4 (Scale Invariance). *The attention mechanism A satisfies:*

$$A(\alpha X) = A(X) \quad \forall \alpha \in \mathbb{R}^+$$

Proof. Consider the attention computation for scaled input:

$$\begin{aligned}A(\alpha X) &= \text{softmax}\left(\frac{(\alpha X)W_QW_K^T(\alpha X)^T}{\sqrt{d}}\right)(\alpha X)W_V \\ &= \text{softmax}\left(\alpha^2 \frac{XW_QW_K^TX^T}{\sqrt{d}}\right)\alpha XW_V\end{aligned}$$

By the scaling property of softmax:

$$\text{softmax}(\alpha^2 Z) = \text{softmax}(Z)$$

Therefore:

$$A(\alpha X) = \text{softmax}\left(\frac{XW_QW_K^TX^T}{\sqrt{d}}\right)XW_V = A(X)$$

\square

4 Natural Gradient Properties

Conjecture 5 (Natural Gradient Convergence). For natural attention with Fisher approximation $F_{NA} = AA^T$, assuming \mathcal{L} is μ -strongly convex in the natural metric:

$$\|\theta_t - \theta^*\|_{F_{NA}}^2 \leq (1 - \eta\mu)^t \|\theta_0 - \theta^*\|_{F_{NA}}^2$$

Remark 6. To prove this conjecture, we need to establish:

1. F_{NA} is positive definite
2. Strong convexity in the natural metric
3. Lipschitz conditions on the gradients

5 Information Processing

Theorem 7 (Basic Information Bound). *For natural attention A and input X :*

$$I(X; A(X)) \leq \log(L)$$

where L is the sequence length.

Proof. By definition of mutual information:

$$I(X; A(X)) = H(A(X)) - H(A(X)|X)$$

Since $A(X)$ is a stochastic matrix (due to softmax):

$$H(A(X)) \leq \log(L)$$

And:

$$H(A(X)|X) \geq 0$$

Therefore:

$$I(X; A(X)) \leq \log(L)$$

□

Conjecture 8 (Stronger Information Bound). Under natural attention optimization:

$$I(X; A(X)) \leq \frac{1}{2\lambda} \mathbb{E}[\|A\|_F^2]$$

Remark 9. To prove this stronger bound, we need to:

1. Connect Frobenius norm to entropy
2. Establish regularization effects on information flow
3. Prove tightness of the bound

6 Future Directions

Several important directions remain for further investigation:

1. Proving the natural gradient convergence conjecture by establishing necessary conditions on F_{NA}
2. Strengthening the information bounds through tighter analysis of attention entropy
3. Developing a complete theory of phase transitions in natural attention
4. Characterizing the relationship between model size and optimal head count

7 Natural Gradient Properties - Completed Proof

Lemma 10 (Positive Definiteness of F_{NA}). *For natural attention matrix A , the Fisher approximation $F_{NA} = AA^T$ is positive definite with minimum eigenvalue $\lambda_{min} \geq \frac{1}{L}$, where L is sequence length.*

Proof. First, note that A is a stochastic matrix due to softmax, so:

$$\sum_j A_{ij} = 1 \quad \forall i$$

For any vector $v \neq 0$:

$$\begin{aligned} v^T F_{NA} v &= v^T A A^T v = \|A^T v\|^2 \\ &= \sum_j \left(\sum_i A_{ij} v_i \right)^2 \\ &\geq \frac{1}{L} \sum_i v_i^2 = \frac{1}{L} \|v\|^2 \end{aligned}$$

The inequality follows from Jensen's inequality and the stochastic property. \square

Lemma 11 (Lipschitz Gradient in Natural Metric). *The loss gradient in natural metric satisfies:*

$$\|\nabla \mathcal{L}(\theta_1) - \nabla \mathcal{L}(\theta_2)\|_{F_{NA}^{-1}} \leq L \|\theta_1 - \theta_2\|_{F_{NA}}$$

Proof. In the natural metric:

$$\|\nabla \mathcal{L}(\theta_1) - \nabla \mathcal{L}(\theta_2)\|_{F_{NA}^{-1}} = \|(F_{NA}^{-1/2}(\nabla \mathcal{L}(\theta_1) - \nabla \mathcal{L}(\theta_2)))\|$$

By softmax properties:

$$\|\nabla \mathcal{L}(\theta_1) - \nabla \mathcal{L}(\theta_2)\| \leq L \|F_{NA}^{1/2}(\theta_1 - \theta_2)\|$$

Combining these gives the result. \square

Theorem 12 (Natural Gradient Convergence - Completed). *For natural attention with Fisher approximation $F_{NA} = AA^T$, the parameter updates satisfy:*

$$\|\theta_t - \theta^*\|_{F_{NA}}^2 \leq (1 - \frac{\eta}{L})^t \|\theta_0 - \theta^*\|_{F_{NA}}^2$$

where η is the learning rate and L is sequence length.

Proof. By the update rule:

$$\theta_{t+1} = \theta_t - \eta F_{NA}^{-1} \nabla \mathcal{L}(\theta_t)$$

In the natural metric:

$$\begin{aligned} \|\theta_{t+1} - \theta^*\|_{F_{NA}}^2 &= \|\theta_t - \theta^* - \eta F_{NA}^{-1} \nabla \mathcal{L}(\theta_t)\|_{F_{NA}}^2 \\ &= \|\theta_t - \theta^*\|_{F_{NA}}^2 - 2\eta \langle \nabla \mathcal{L}(\theta_t), \theta_t - \theta^* \rangle + \eta^2 \|\nabla \mathcal{L}(\theta_t)\|_{F_{NA}^{-1}}^2 \end{aligned}$$

By strong convexity in natural metric (from Lemma 1):

$$\langle \nabla \mathcal{L}(\theta_t), \theta_t - \theta^* \rangle \geq \frac{1}{L} \|\theta_t - \theta^*\|_{F_{NA}}^2$$

By Lipschitz property (from Lemma 2):

$$\|\nabla \mathcal{L}(\theta_t)\|_{F_{NA}^{-1}}^2 \leq L \|\theta_t - \theta^*\|_{F_{NA}}^2$$

Combining:

$$\|\theta_{t+1} - \theta^*\|_{F_{NA}}^2 \leq (1 - \frac{2\eta}{L} + \frac{\eta^2}{L}) \|\theta_t - \theta^*\|_{F_{NA}}^2$$

For $\eta \leq 1$:

$$1 - \frac{2\eta}{L} + \frac{\eta^2}{L} \leq 1 - \frac{\eta}{L}$$

Iterating completes the proof. \square

8 Information Processing - Completed Proof

Lemma 13 (Entropy Bound for Attention). *For attention matrix A :*

$$H(A) \leq \frac{1}{2} \|A\|_F^2 + \log L$$

Proof. By definition of entropy and Jensen's inequality:

$$\begin{aligned} H(A) &= - \sum_{i,j} A_{ij} \log A_{ij} \\ &\leq \frac{1}{2} \sum_{i,j} A_{ij}^2 + \log(\sum_{i,j} 1) \\ &= \frac{1}{2} \|A\|_F^2 + \log L^2 \end{aligned}$$

\square

Theorem 14 (Strong Information Bound). *Under natural attention optimization with regularization parameter λ :*

$$I(X; A(X)) \leq \frac{1}{2\lambda} \mathbb{E}[\|A\|_F^2] + \log L$$

Proof. The mutual information decomposes as:

$$I(X; A(X)) = H(A(X)) - H(A(X)|X)$$

By Lemma on Entropy Bound:

$$H(A(X)) \leq \frac{1}{2} \|A\|_F^2 + \log L$$

The regularization term in optimization:

$$\mathcal{L}_{reg} = \lambda \|A\|_F^2$$

implies:

$$\mathbb{E}[\|A\|_F^2] \leq \frac{1}{\lambda} \mathbb{E}[\mathcal{L}]$$

Since $H(A(X)|X) \geq 0$:

$$I(X; A(X)) \leq \frac{1}{2\lambda} \mathbb{E}[\|A\|_F^2] + \log L$$

□

9 Future Directions

Several important directions remain for further investigation:

1. Strengthening the information bounds through tighter analysis of attention entropy
2. Developing a complete theory of phase transitions in natural attention
3. Characterizing the relationship between model size and optimal head count